

VARIABLE PRECISION DEEP LEARNING (FLOAT16 & FLOAT32)

- **CS-597: Reading and Special Problems**
- **Team Members**
 - Nilesh Jorwar (A20405042)
 - Ashish Ryot (A20405230)

PROBLEM STATEMENT

- Performance of CPU with low precision data types, along with data representation techniques (floating point) and the variation of results with adjustments to various parameters of the model needs to be experimented.
- Improvisation of the previous project to get better accuracy and lower training time using various techniques:
 - Implementation of AVX instructions for speeding up runtime (float32)
 - Multithreading implementation for float32 and float16 data types

MOTIVATION


- Need for speedup and performance on larger datasets using low variable precision data types without loss of accuracy, precision and recall.
- Deep learning on problems like image recognition is very compute intensive task.

IMPROVEMENT OVER PAST PROJECT & CURRENT PROGRESS

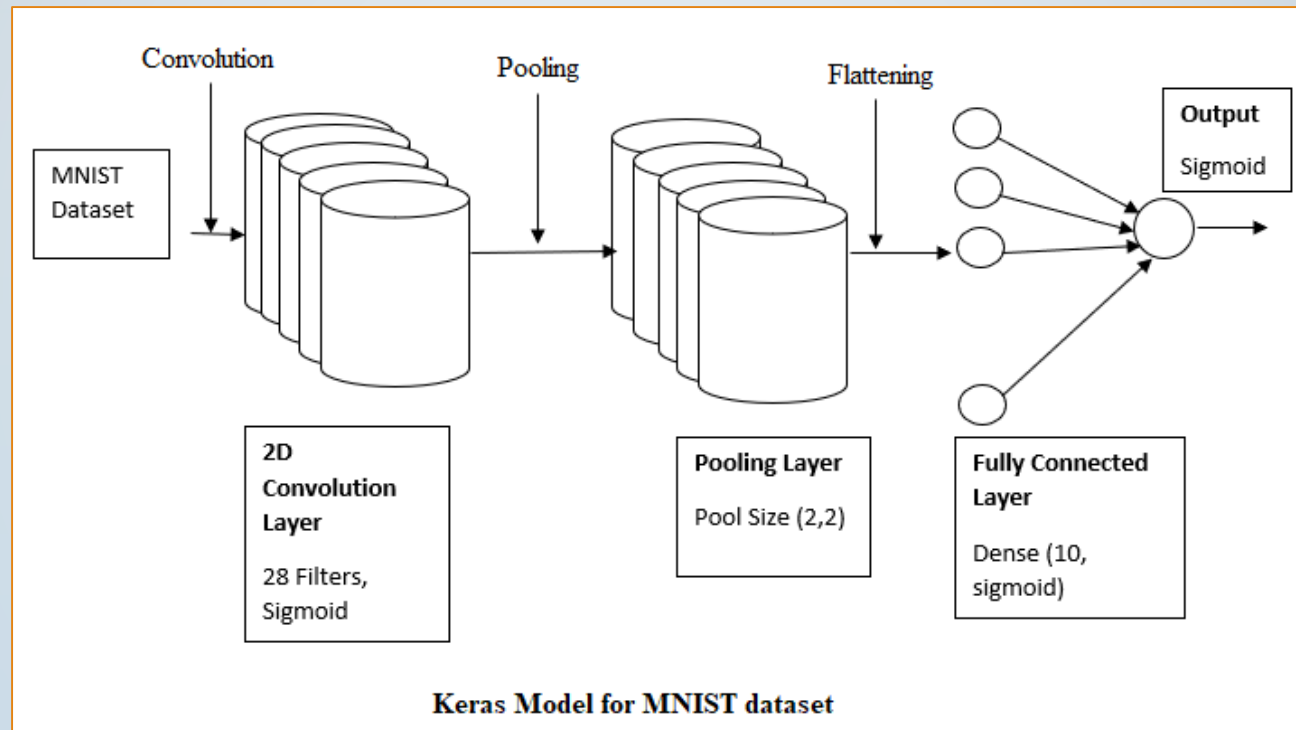
- AVX instructions for float32 data types
- Parallelization of neural network using thread library for float16 data types
- Testing model on more complex datasets
- Comparison of C++ model with TensorFlow and Keras in Python
- Testing of Keras model with int data type

ARCHITECTURES

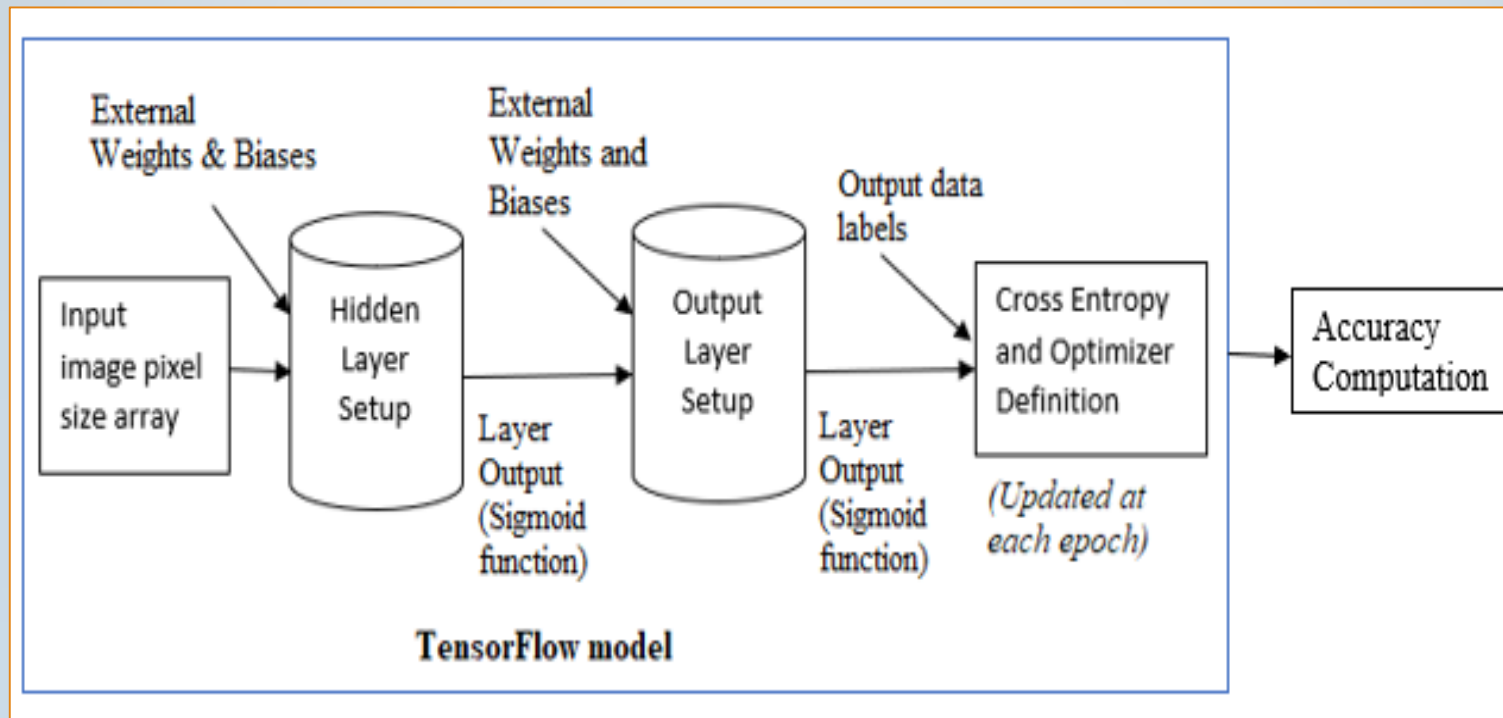
- cpp Model – Input layer with 784 neurons, hidden layer (variable neurons), and output layer (10 neurons)
 - Activation Functions: sigmoid
- TensorFlow Python Model
- Keras Python Model



ARCHITECTURE – Keras Model



ARCHITECTURE – TensorFlow Model



TESTBED REQUIREMENTS

- Intel(R) Xeon(R) CPU e5-2667 v3 @ 3.20ghz
- CPU 2
- Cores 16
- Logical Processors 32

ACHIEVEMENTS

- Successful implementation of multithreading on neural network model
- AVX instructions implementation
- Improved accuracy, precision and recall
- Reduced runtime.
- Experiments over different datasets
- Evaluation over different platforms and models like Keras, TensorFlow, and CPP

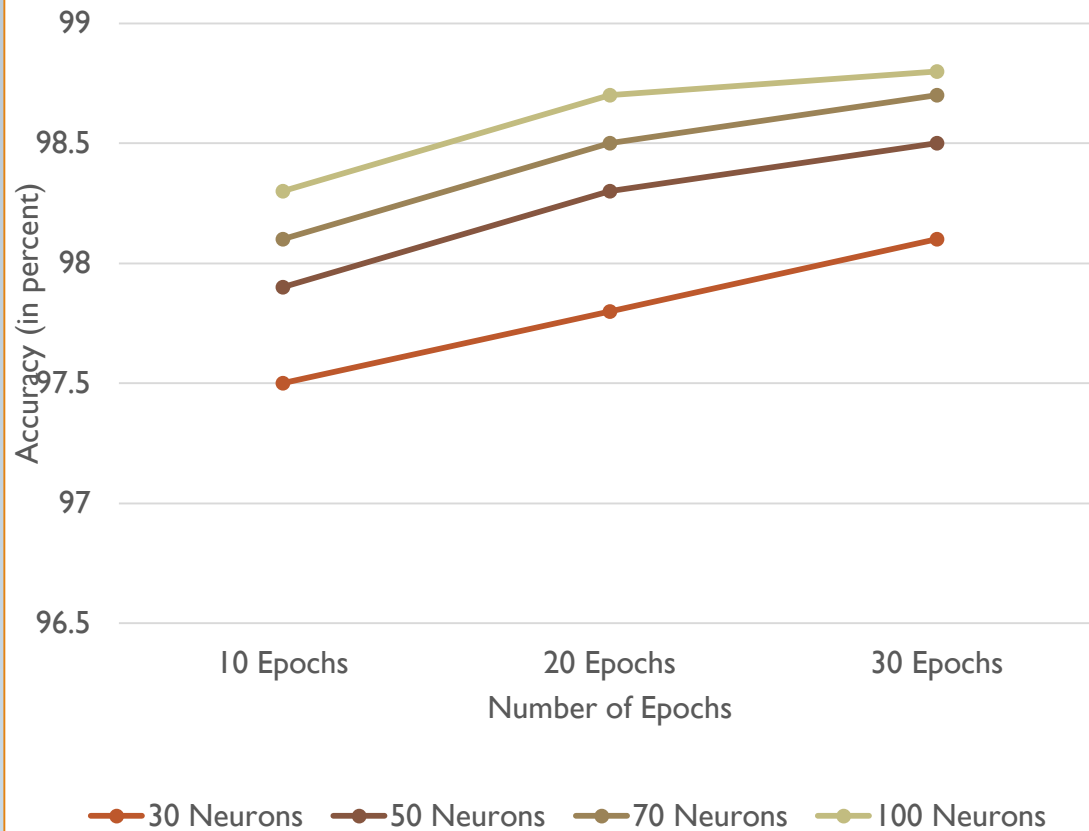
CHALLENGES FACED

- Implementation of AVX512 instructions for float16 data types.
- Implementation using int data types.
- Implementation using low precision data types (CLN) [4].
 - The external library used, rounds the original data into fixed number of real bits , which results in loss of precision.
 - The precision decreases and number of epochs required to get good accuracy increases.

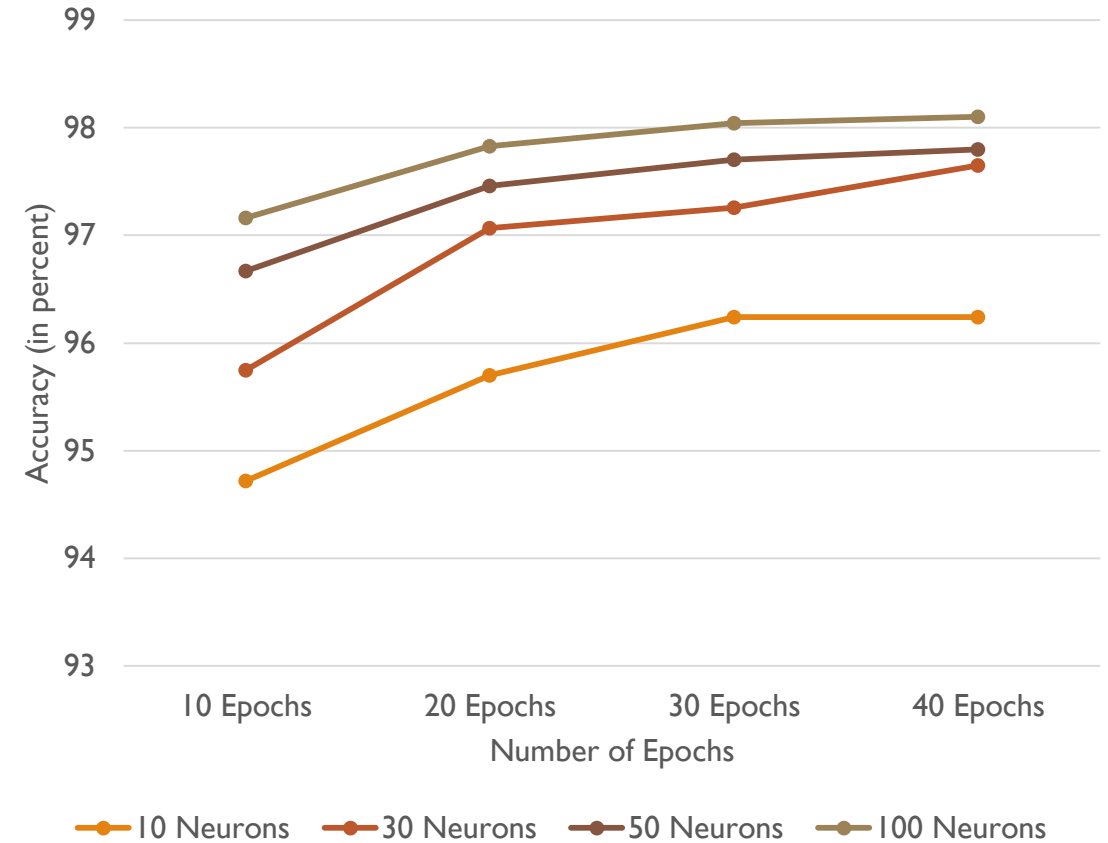
PERFORMANCE EVALUATION

- MNIST and Fashion dataset
 - Float32 and float16 data types
- Evaluation using varying number of nodes like 10, 20, 30, 50, 70, 100, and 400 in hidden layer
- Accuracy computations for different number of epochs like 10, 20, 30, and 35.

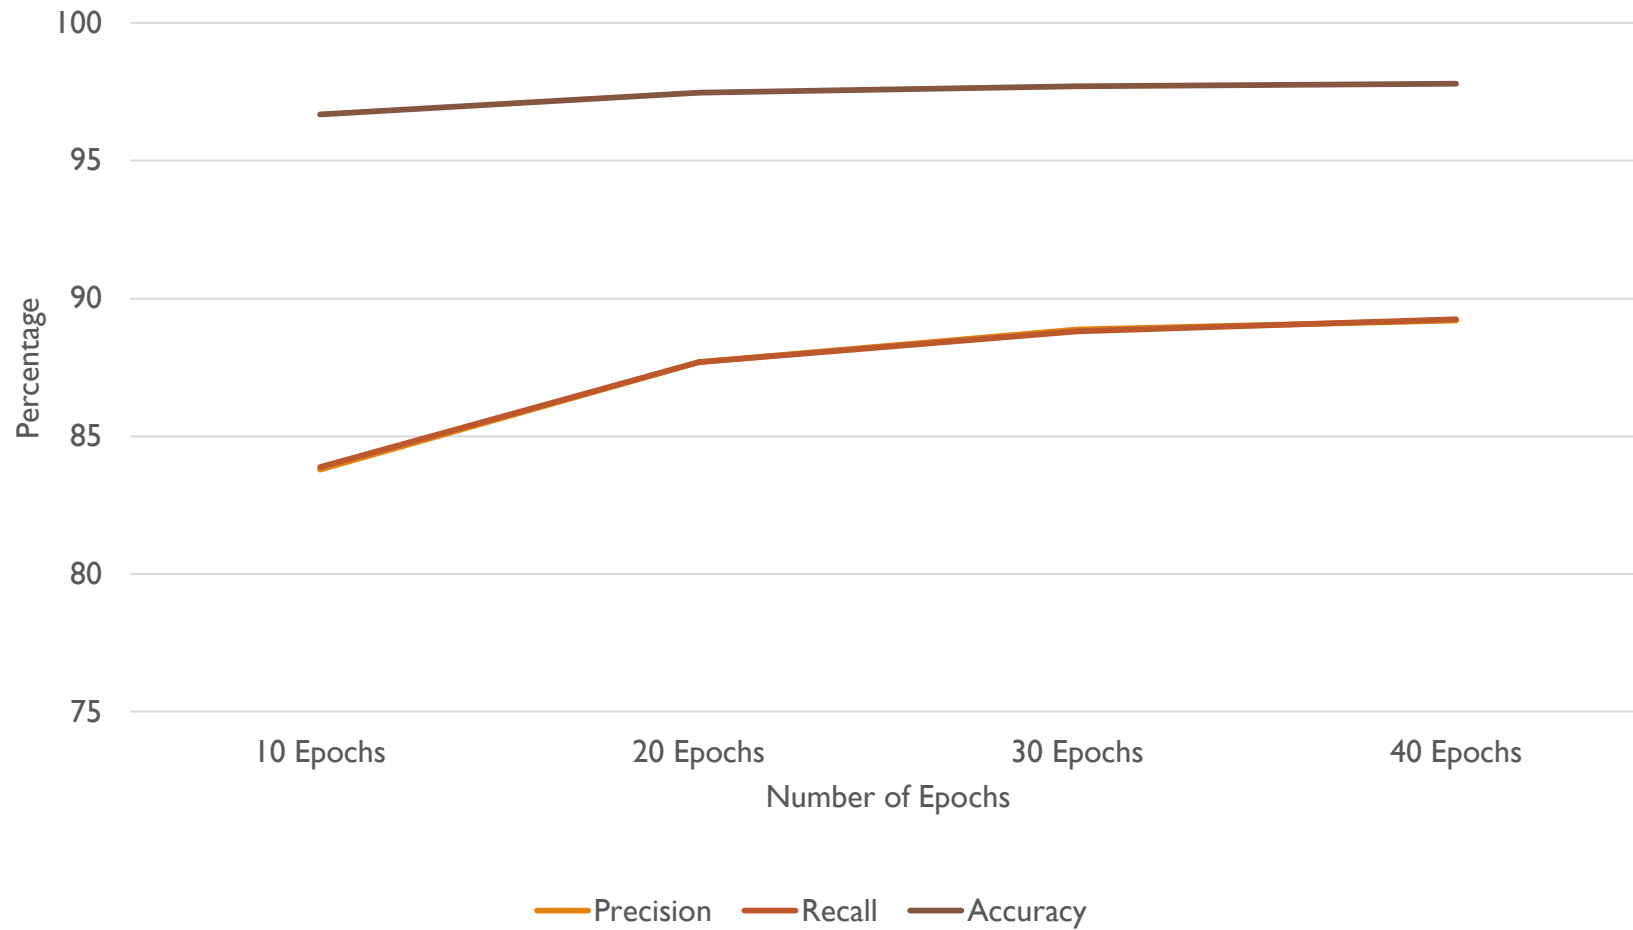
Number of Epochs vs Accuracy for different number of neurons in hidden layer (Float 32)



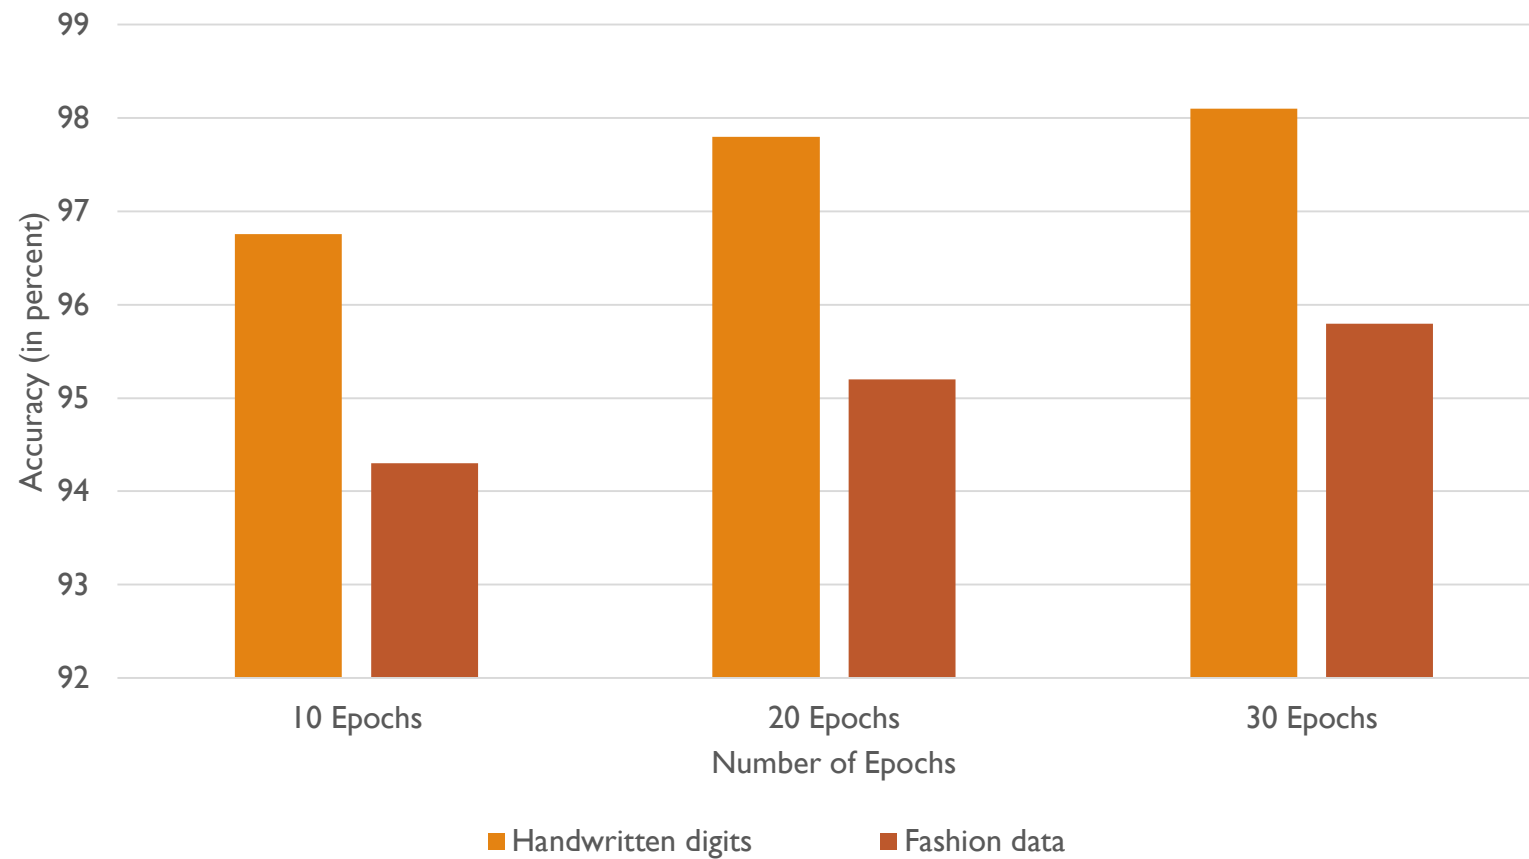
Number of Epochs vs Accuracy for different number of neurons in hidden layer (Float 16)



Precision, Recall and Accuracy for MNIST dataset



Comparison of model on MNIST Handwritten dataset vs Fashion dataset




Accuracy and Runtime over different epochs for different models

		cpp Model Float32 (avx)	cpp Model Float16 (multi- threaded)	TensorFlow Model (Keras) Float16
10 Epochs	Accuracy (%)	96.8	95.72	98.32
	Runtime (secs)	40.51	654	761
20 Epochs	Accuracy (%)	97.8	97.07	98.43
	Runtime (secs)	133.50	1593	1530
30 Epochs	Accuracy (%)	98.1	97.26	98.55
	Runtime (secs)	205.47	2425	2276

		cpp Model Float32	TensorFlow Model Float32
10 Epochs	Accuracy (%)	96.2	95.69
	Runtime (secs)	51.33	30
20 Epochs	Accuracy (%)	97.8	95.65
	Runtime (secs)	152.52	60.928
30 Epochs	Accuracy (%)	97.9	95.72
	Runtime (secs)	231.37	90.88

Training Error per Epoch for different models



FUTURE SCOPE

- Implementation of VNNI instructions on float16 data types
- Conversion of CPP model to work on GPU platform using CUDA
- Testing the CPP model on ImageNet

REFERENCES

- [1] <https://thinkingandcomputing.com/posts/using-avx-instructions-in-matrix-multiplication.html>
- [2] <https://adventuresinmachinelearning.com/python-tensorflow-tutorial/>
- [3] <https://cognitivedemons.wordpress.com/2017/07/06/a-neural-network-in-10-lines-of-c-code/>
- [4] <https://ginac.de/CLN/cln.html>